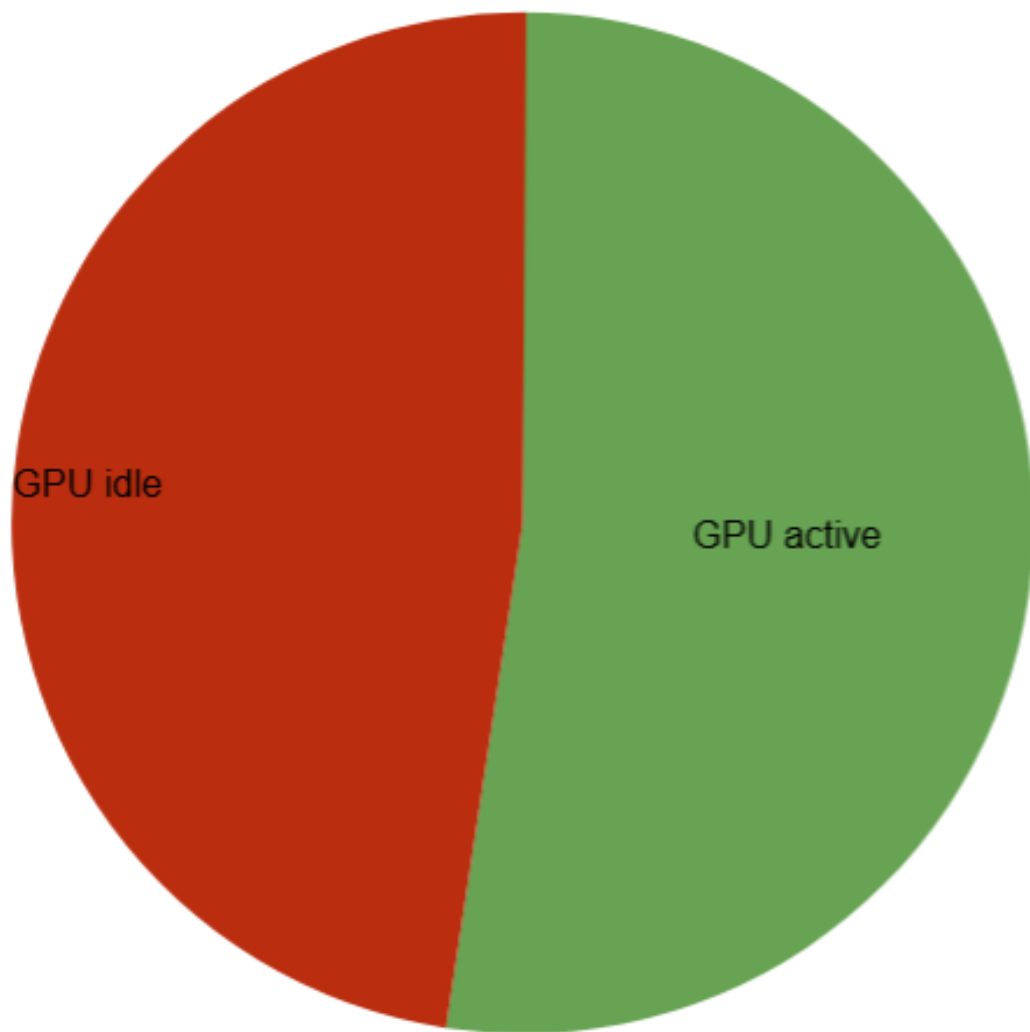



GPU Active Time ?



Legend

- GPU active
- GPU idle

GPU Utilization Over Time ?



legend

— algo-1 — algo-2

Group utilization by Host ▼

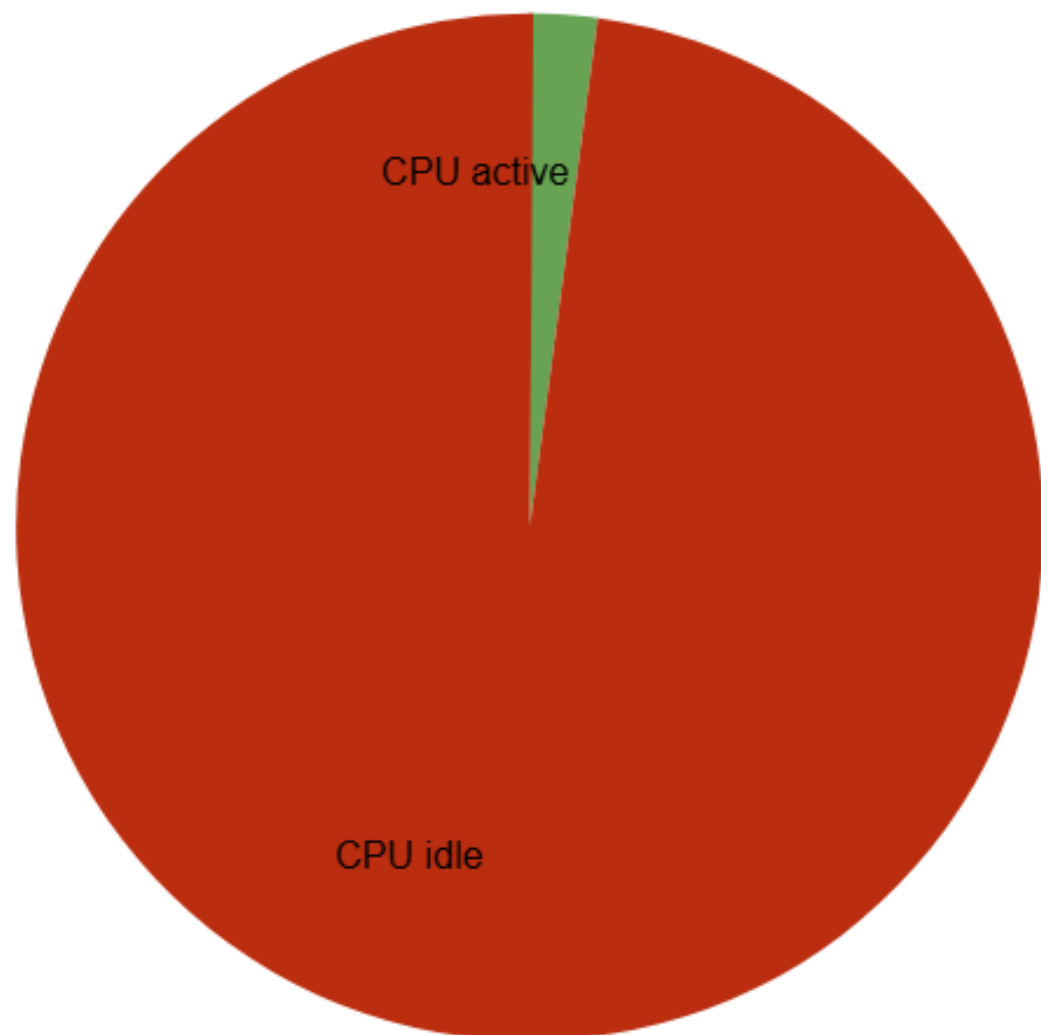
Open in Timeline

Reset to default


Zoom into selected range

Set current range for CPU

CPU Active Time ?

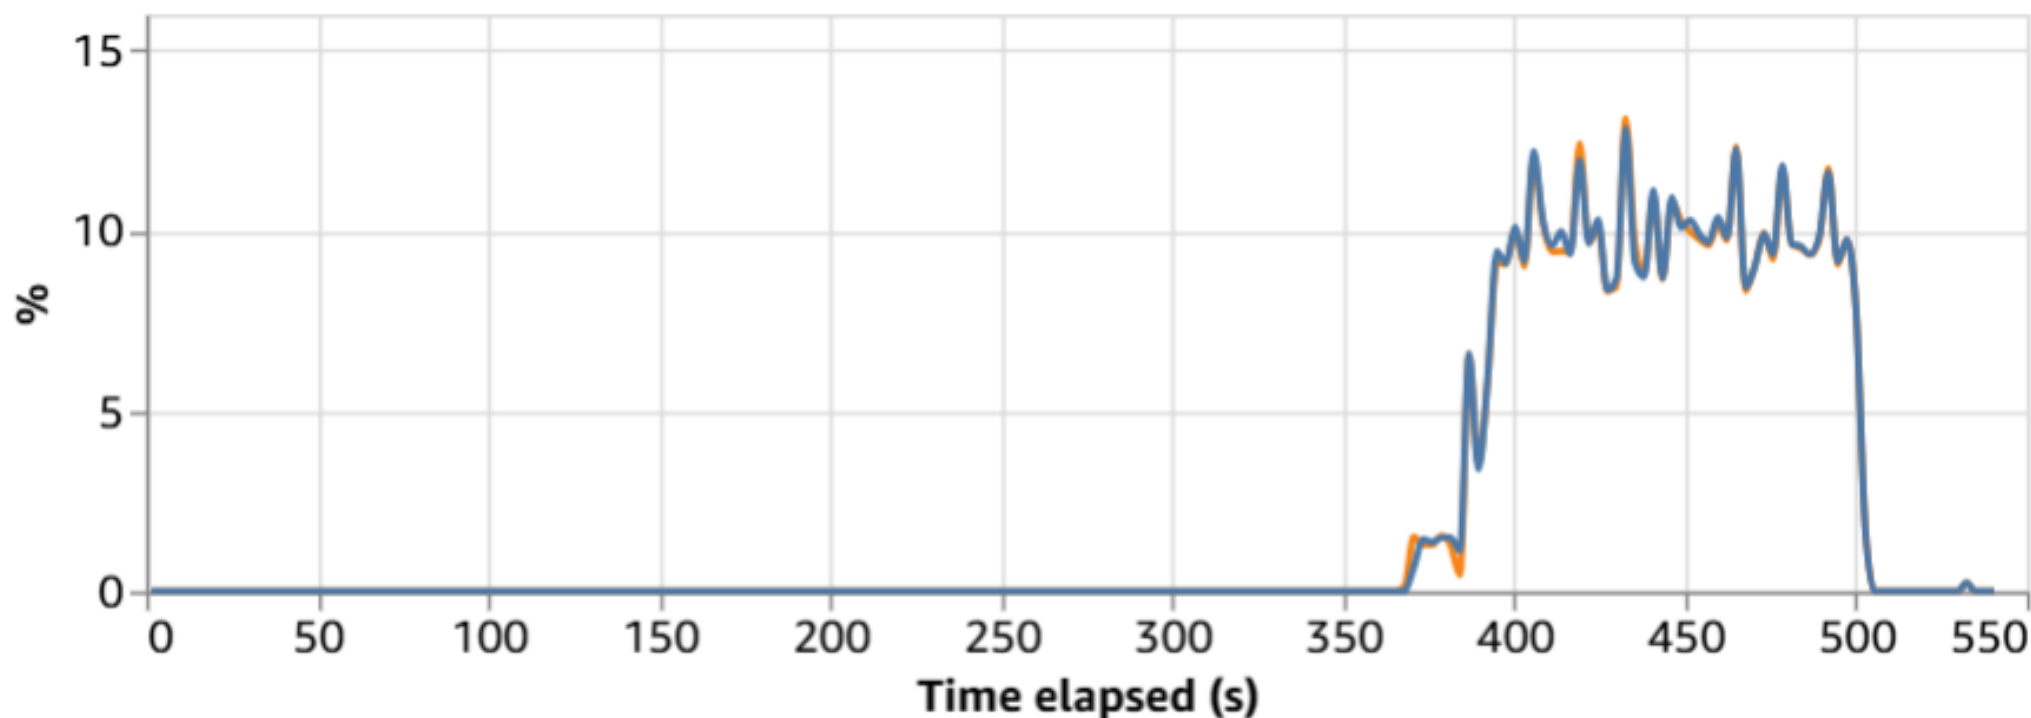


Legend

 CPU active

 CPU idle

CPU Utilization Over Time ?



legend


— algo-1 — algo-2

Open in Timeline

Reset to default


Zoom into selected range

Time Spent By All GPU Kernels ?



k 15

Time Spent By Top 15 GPU Kernels ?




Legend

- ncclKernel_AllReduce_RI...
- ncclKernel_AllReduce_T...
- void at::native::(anonym...
- void at::native::unrolled_...
- void cudnn::ops::nchwTo...
- void at::native::(anonym...
- void at::native::(anonym...
- void at::native::(anonym...
- DtoD
- sm75_xmma_fprop_imp...
- void at::native::vectorize...
- void at::native::(anonym...
- void at::native::(anonym...
- void at::native::(anonym...
- void at::native::(anonym...

Launch Counts Of All GPU Kernels ?

...




Legend

- void cudnn::ops::nchwTo...
- void at::native::vectorize...
- DtoD
- void at::native::unrolled_...
- void at::native::vectorize...
- void cudnn::ops::nhwcTo...
- void cudnn::bn_fw_inf_1...
- void at::native::unrolled_...
- void at::native::vectorize...
- void at::native::vectorize...
- sm75_xmma_fprop_imp...
- Memset
- void at::native::vectorize...
- void cudnn::bn_fw_tr_1C...
- void cudnn::bn_bw_1C1...
- Others

k 15

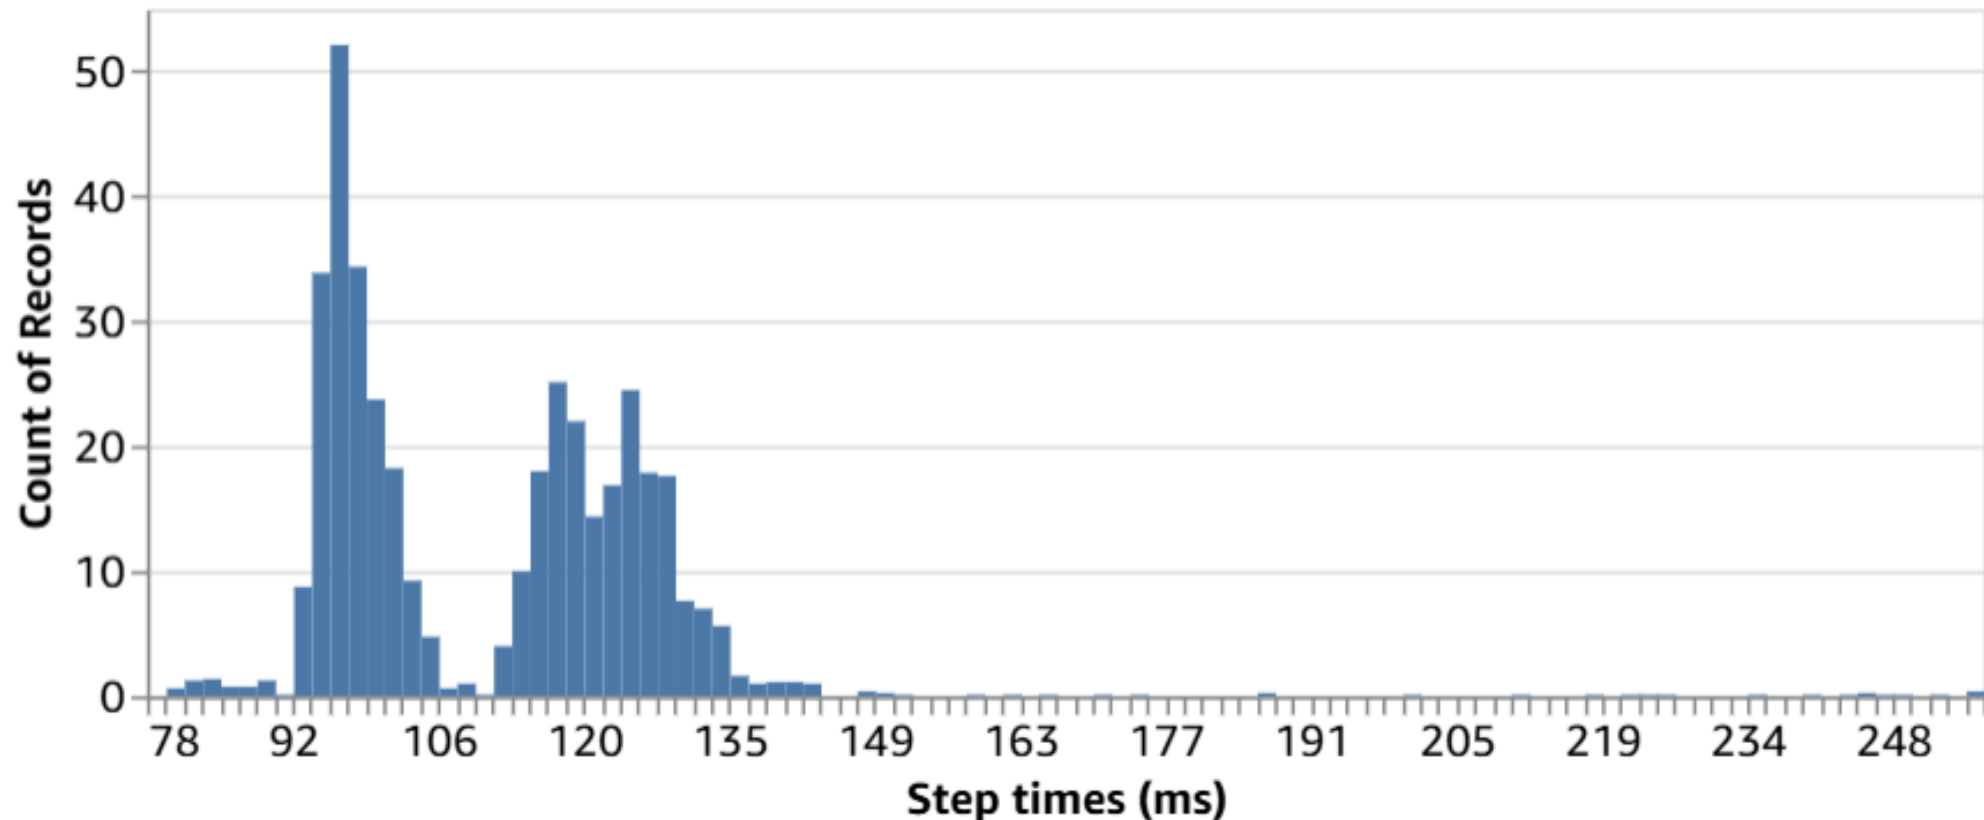
Launch Counts Of Top 15 GPU Kernels ?



Legend

- void cudnn::ops::nchwTo...
- void at::native::vectorize...
- DtoD
- void at::native::unrolled_...
- void at::native::vectorize...
- void cudnn::ops::nhwcTo...
- void cudnn::bn_fw_inf_1...
- void at::native::unrolled_...
- void at::native::vectorize...
- void at::native::vectorize...
- sm75_xmma_fprop_imp...
- Memset
- void at::native::vectorize...
- void cudnn::bn_fw_tr_1C...
- void cudnn::bn_bw_1C1...

Step Time Distribution ?



nBins 100

xStart 79

xEnd 256

☒ Remove outliers

☐ logAxis

Mean 111.23 ms


Integral 3146

Standard Deviation 16.79 ms


Underflow count 0

Overflow count 62

Kernel Precision Distribution ?



GPU Activity Distribution ?




Legend

- kernel
- memcpy
- memset
- sync

GPU Memory Operations Distribution ?

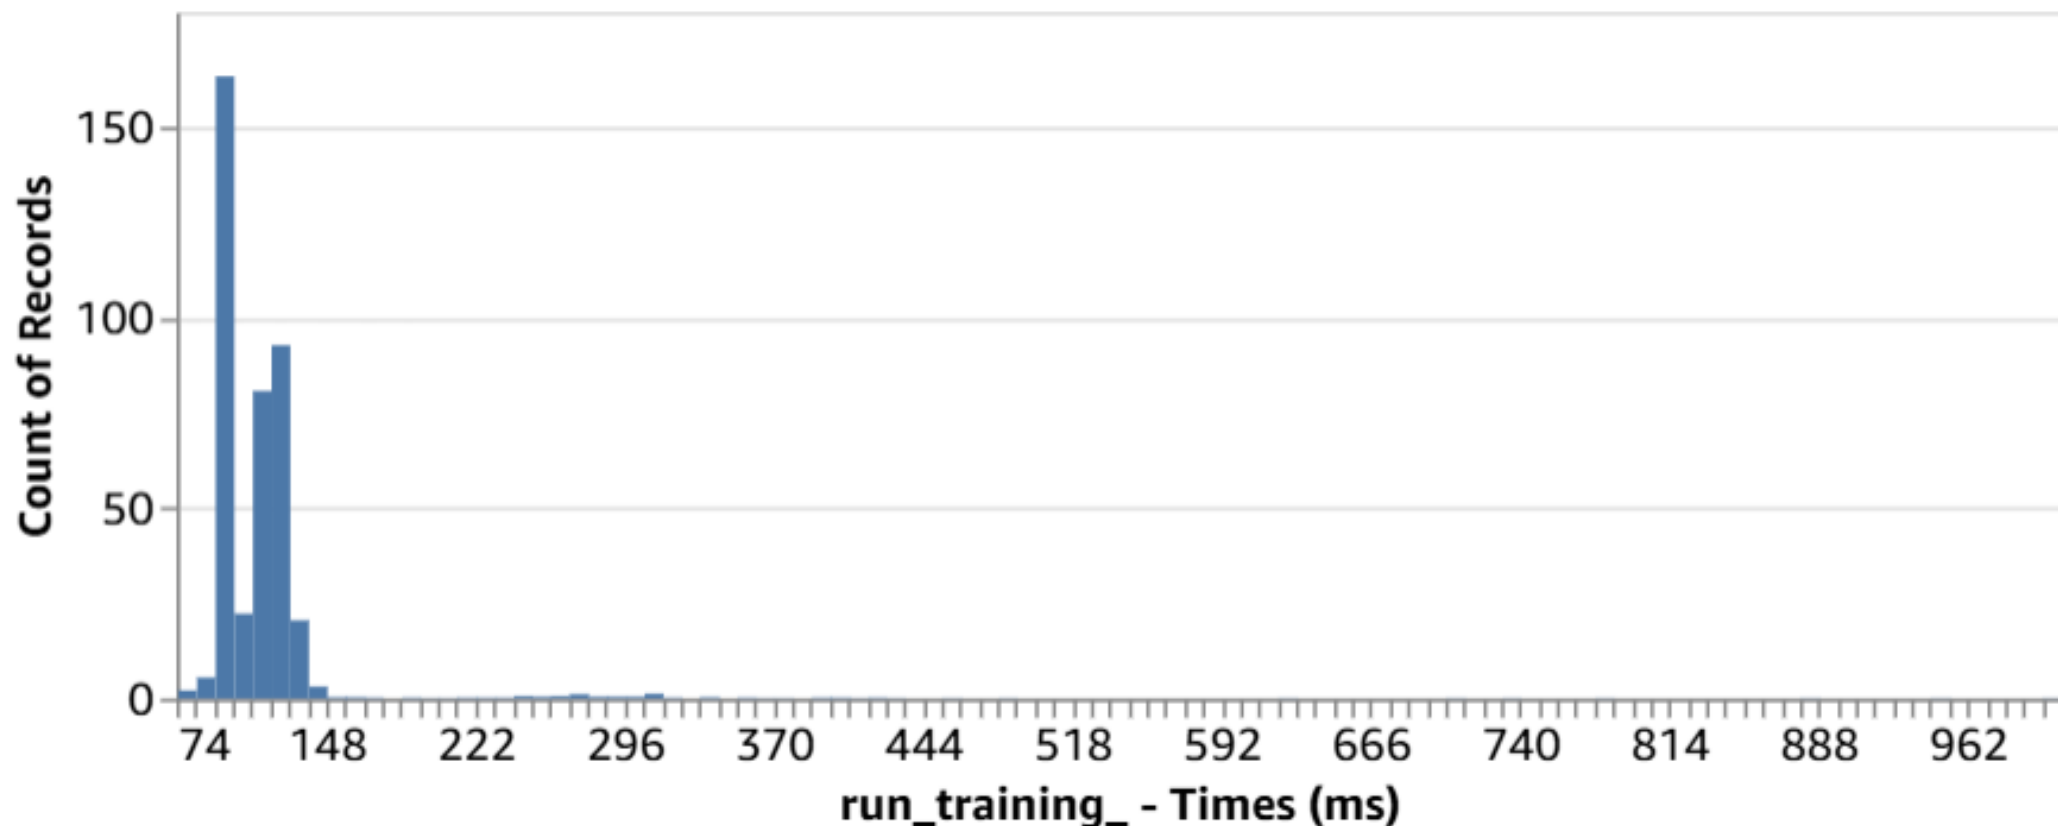
...



Legend

- DtoD.device.device
- HtoD.pageable.device
- DtoH.device.pageable
- HtoD.pinned.device

Time distribution for: run_training



nBins 100

xStart 79

xEnd 1004

Mean 116.12 ms

Integral 3208

Standard Deviation 46.53 ms

Underflow count 0

Overflow count 40

☒ Remove outliers

☐ logAxis